

25 Spring ECEN 610: Mixed-Signal Interfaces

Lab2: Signal to Noise Ratio, Quantization

Name: Yu-Hao Chen

UIN:435009528

Section:601


Professor: Sebastian Hoyos

TA: Sky Zhao

GitHub: https://github.com/Yu-HaoChen/TAMU_ECEN610_Mixed_signal/tree/main

1. SIGNAL TO NOISE RATIO (40%) Generate a tone with frequency 2 MHz and amplitude 1 V. Sample the tone at frequencies $F_s = 5$ MHz.

a) Add Gaussian noise to the sampled sinewave such that the signal SNR is 50 dB. Find first the variance of the Gaussian noise needed to produce the target SNR. Calculate and plot the Power Spectral Density (PSD) from the DFT of the noisy samples. Corroborate that the SNR calculation from the DFT plot gives the theoretical result. What would be the variance of a uniformly distributed noise to obtain the same SNR.



$$X(t) = \cos(2\pi f_0 t) \quad f_s = 5M \quad \text{nyquist} = 2.5M \quad \text{SNR} = 50\text{dB}$$

$$f_0 = 2M$$


$$\text{Set } N_{\text{DFT}} = 64 \quad \Delta f = \frac{5M}{64} = 78.125 \text{ kHz} \quad \text{bin range}$$


$$\text{SNR} = 10 \log_{10} \frac{P_{\text{signal}}}{P_{\text{noise}}} \quad \begin{array}{c} 78.125 \text{ kHz} \\ 0 \quad 1 \quad 63 \end{array}$$

$$P_{\text{signal}} = \frac{1}{T} \int_0^T x^2(t) dt = \frac{A^2}{2} \quad \left. \begin{array}{l} \text{continuous} \\ \text{time domain} \\ \text{discrete} \end{array} \right\} x(n) = \cos$$


$$= \frac{1}{N} \sum_{n=0}^{N-1} x[n]^2$$

$$= \frac{1}{N} \sum_{k=0}^{N-1} S_{xx}[k]^2 \quad \left. \begin{array}{l} S_{xx} \text{ power spectral density} \\ \frac{x[k]^2}{N} \end{array} \right\} \text{freq domain}$$







b) Now repeat a.) using a window before the DFT. Use the following windows: Hanning, Hamming, Blackman. What are your conclusions? NOTE: The use of windows mentioned above spreads the signal power. You must take this into account when computing SNR.



Hanning window




Hamming window




Blackman window

2. QUANTIZATION

- a) Create a perfect quantizer with 6 bits of resolution and with flexible sampling rate. For a 200 MHz full scale input tone, sample and quantize the sinewave at 400 MHz and plot the PSD of 30 periods. What is the SNR? Repeat the SNR calculation for 100 periods of the same signal. Make your own conclusions about this test regarding periodicity of quantization noise and the impact of this in the SNR. How can you solve this problem?



```
=== SNR for 30 Cycles ===
SNR = -240.12 dB
```




```
=== SNR for 100 Cycles ===
SNR = -229.76 dB
```

- The periodicity of quantization noise becomes most pronounced when the input frequency f_{in} and the sampling rate f_s form a simple integer ratio. This can cause significant variations in the measured SNR depending on the number of sampled cycles, the record length, and how the FFT bins are aligned.
- By breaking this periodicity—such as by slightly adjusting the input frequency, adding dither, or using an appropriate window and record length design—the quantization noise can be spread across the entire bandwidth. This makes the noise more closely resemble ideal “white noise,” and thus yields an SNR measurement that is closer to the theoretical value.

b) Find an incommensurate sampling frequency larger than Nyquist rate. Plot the PSD

of the new samples. Calculate the SNR from the figure.




=== SNR for 30 Cycles ===
SNR = 35.68 dB

c) Repeat a) using a 12 bit quantizer. Can you prove from simulations that $\text{SNR} \sim 6N$ (where N is the number of bits used by the quantizer) in both the cases, $N = 6$ and $N = 12$?


$$\text{SNR} = \frac{P_{\text{signal}}}{P_{\text{noise}}} = \frac{\frac{V_{\text{ref}}^2}{2}}{\frac{\Delta V^2}{12}} = \frac{\frac{V_{\text{ref}}^2}{2}}{\frac{2^{2N} V_{\text{ref}}^2}{12}} = \frac{\frac{V_{\text{ref}}^2}{2}}{\frac{4 V_{\text{ref}}^2}{12 \cdot 2^{2N}}} = \frac{6 \cdot 2^{2N}}{4} = \frac{3}{2} \cdot 2^{2N}$$

$$10 \log_{10} \frac{3}{2} \cdot 2^{2N} = 10 \log_{10} \frac{3}{2} + 10 \cdot 2N \log_{10} 2 = 6N + 1.72$$



=== SNR for 30 Cycles ===
SNR = 74.89 dB


d) Use a Hanning window and repeat c). What is the SNR? Make your own conclusions.



```
=== SNR for 30 Cycles ===
SNR = 74.89 dB
```

The window will have minor effect on intergrate cycle casue there are no specturm leakage for intergrate cycle

e) Now add noise again so the signal SNR is 38 dB . Repeat c) and d). What are the SNRs? Provide conclusions.



```
=== Measured SNR ===
SNR = 34.33 dB
p_quant_error = 1.621523e-08
p_noise_total = 1.585000e-04
p_noise       = 1.584838e-04 ( = p_noise_total - p_quant_error )
```

Appendix

Q1.a PSD noise (without quantization)


```

10 # === continuous time signal ===
11 t_cont = np.linspace(0, N * Ts, 1000) # Continuous time axis (1000 points)
12 x_t = np.cos(2 * np.pi * F * t_cont) # Continuous time signal x(t)
13
14 # === discrete time signal ===
15 t_n = np.arange(N) * Ts # Discrete time axis: 64 points
16 x_n = np.cos(2 * np.pi * F * t_n) # Sampled signal x[n]
17
18 # === POWER signal & noise ===
19 signal_power = np.mean(x_n ** 2) # P signal= 1/N (Σ (n=0~N-1) x[n]²), x[n] time domain
20 noise_power = signal_power / (10 ** (SNR_dB / 10)) # SNR= 10 log10(P signal/ P noise)
21 noise = np.sqrt(noise_power) # avg= 0 white noise, P noise= σ² noise
22
23 # === Gaussian Noise ===
24 Gaussian_noise = np.random.normal(0, noise, N) # Gaussian noise
25 x_noisy = x_n + Gaussian_noise # Noisy signal
26
27 # === FFT ===
28 X_k = np.fft.fft(x_noisy, N) # DFT freq amp x[k]
29 X_k_shifted = np.fft.fftshift(X_k) # Move center to 0
30 frequencies = np.fft.fftfreq(N, Ts) # DFT freq (0, 78.125k, ..., 2.5M)
31 frequencies_shifted = np.fft.fftshift(frequencies) # Move center to 0
32
33 # === PSD ===
34 df = Fs / N # bin value: 78.125k
35 PSD = (np.abs(X_k_shifted) ** 2 / (N * Fs)) * df # [PSD(bin)/fs= power/hz] * bin's Hz (df)= PSD power/bin
36 PSD_dB = 10 * np.log10(PSD) # dB

```

Hanning window

```

4 Fs = 5e6
5 F = 2e6
6 N = 64
7 Ts = 1 / Fs
8 SNR_dB = 50
9
10 # === 1.(Sampled Signal) ===
11 t_n = np.arange(N) * Ts # discrete time axis
12 x_n = np.cos(2 * np.pi * F * t_n) # sample signal
13
14 # === 2.Gaussian noise ===
15 signal_power = np.mean(x_n ** 2) # signal power
16 noise_power = signal_power / (10 ** (SNR_dB / 10)) # noise power
17 noise_std = np.sqrt(noise_power) # noise standard
18 Gaussian_noise = np.random.normal(0, noise_std, N) # Gaussian Noise
19 x_noisy = x_n + Gaussian_noise # signal with noise
20
21 # === 3. Hanning Window ===
22 hanning_window = np.hanning(N) # Hanning Window
23 x_windowed = x_noisy * hanning_window # signal with noise + Window
24
25 # === 4. FFT & freq axis ===
26 X_k = np.fft.fft(x_windowed, N) # FFT Spectrum
27 X_k_shifted = np.fft.fftshift(X_k) # center 0
28 frequencies = np.fft.fftfreq(N, Ts) # FFT freq axis (0, 78.125k, ... 2.5M)
29 frequencies_shifted = np.fft.fftshift(frequencies) # center 0
30
31 # === 5. Power Spectral Density (PSD) ===
32 df = Fs / N
33 PSD = (np.abs(X_k_shifted) ** 2 / (N * Fs)) * df # PSD
34 PSD_dB = 10 * np.log10(PSD) # dB 單位

```

Hamming window

```

21 # === 3. Hamming Window ===
22 hanning_window = np.hamming(N) # Hanning Window
23 x_windowed = x_noisy * hanning_window # signal with noise + Window

```

Blackman window

```

21 # === 3. Blackman Window ===
22 hanning_window = np.blackman(N) # Hanning Window
23 x_windowed = x_noisy * hanning_window # signal with noise + Window

```

Q2


```

10 def simulate_quantization(N_cycles):
11     # === 1. x[n] ===
12     N = int(N_cycles * Fs / F) # N= int (period) * (T0/Ts= sample point in one T0)
13     t_n = np.arange(N) / Fs     # x[n] time axis: n * Ts, n= 0, 1, 2, 3... N-1
14     t_cont = np.linspace(0, N / Fs, 1000) # x(t) time axis
15     x_input_cont = Vref * np.sin(2 * np.pi * F * t_cont) # x(t)
16     x_input = Vref * np.sin(2 * np.pi * F * t_n) # x[n]
17
18     # === 2. Quantization ===
19     x_quantized_con = np.round((x_input_cont + Vref) * (quantization_levels - 1) / (2 * Vref)) * (2 * Vref) / (quantization_levels - 1) - Vref
20     x_quantized = np.round((x_input + Vref) * (quantization_levels - 1) / (2 * Vref)) * (2 * Vref) / (quantization_levels - 1) - Vref
21     x_quantized_nodc = x_quantized - np.mean(x_quantized)
22     quantization_noise = x_input - x_quantized # Quantization error
23     print(x_quantized)
24     quantization_noise_nodc = x_input - x_quantized_nodc
25
26     # === 3. SNR ===
27     signal_power = np.mean(x_input ** 2) # P signal
28     noise_power = np.mean(quantization_noise ** 2)
29     SNR = 10 * np.log10(signal_power / noise_power)

```

Q2 window

```

11 def simulate_quantization(N_cycles):
12     N = int(N_cycles * Fs / F)
13     t_n = np.arange(N) / Fs
14     t_cont = np.linspace(0, N / Fs, 1000)
15
16     x_input_cont = Vref * np.sin(2 * np.pi * F * t_cont)
17     x_input = Vref * np.sin(2 * np.pi * F * t_n)
18
19     x_quantized = np.round((x_input + Vref) * (quantization_levels - 1) / (2 * Vref)) \
20     | * (2 * Vref) / (quantization_levels - 1) - Vref
21     quantization_noise = x_input - x_quantized
22
23
24     signal_power = np.mean(x_input ** 2)
25     noise_power = np.mean(quantization_noise ** 2)
26     SNR = 10 * np.log10(signal_power / noise_power)
27
28     window = np.hanning(N)
29     x_windowed = x_quantized * window
30
31     X_k = np.fft.fft(x_windowed, N)
32     X_k_shifted = np.fft.fftshift(X_k)
33
34     frequencies = np.fft.fftfreq(N, 1 / Fs)
35     frequencies_shifted = np.fft.fftshift(frequencies)
36
37     df = Fs / N
38     PSD = (np.abs(X_k_shifted) ** 2 / (N * Fs)) * df
39     PSD_dB = 10 * np.log10(PSD + 1e-20)

```

Q2 noise

```

10 def simulate_quantization_5plots(N_cycles, p_noise_total=1.585e-4):
11     # 1) Generate discrete-time signal x[n] for N integer cycles
12     N = int(N_cycles * Fs / F)
13     t_n = np.arange(N) / Fs
14     t_cont = np.linspace(0, N / Fs, 1000) # For continuous plotting
15
16     # Create a clean sinusoidal input signal
17     x_input_cont = Vref * np.sin(2 * np.pi * F * t_cont) # continuous
18     x_input = Vref * np.sin(2 * np.pi * F * t_n) # discrete
19
20     # 2) Quantize the clean signal
21     x_quantized = np.round((x_input + Vref) * (quantization_levels - 1) / (2 * Vref)) \
22     | * (2 * Vref) / (quantization_levels - 1) - Vref
23
24     # 3) Compute original quantization error power
25     p_quant_error = np.mean((x_input - x_quantized)**2)
26
27     # 4) Compute p_noise = p_noise_total - p_quant_error
28     p_noise = p_noise_total - p_quant_error
29     if p_noise < 0:
30         p_noise = 0.0
31     sigma = np.sqrt(p_noise)
32
33     # Generate Gaussian noise with variance = p_noise
34     G_noise = np.random.normal(0, sigma, size=x_quantized.shape)
35
36     # Final quantization noise = (x_input - x_quantized) + G_noise
37     noise = (x_input - x_quantized) + G_noise
38     noise_nodc = noise - np.mean(noise)
39
40     # 5) Create a Hann window for the quantized signal
41     window = np.hanning(N)
42     x_windowed = x_quantized * window

```

```
40     # 5) Create a Hann window for the quantized signal
41     window = np.hanning(N)
42     x_windowed = x_quantized * window
43
44     # 6) Compute FFT & PSD (Hann-windowed)
45     X_k = np.fft.fft(x_windowed, N)
46     X_k_shifted = np.fft.fftshift(X_k)
47     freqs = np.fft.fftfreq(N, 1 / Fs)
48     freqs_shifted = np.fft.fftshift(freqs)
49
50     df = Fs / N
51     psd = (np.abs(X_k_shifted) ** 2 / (N * Fs)) * df
52     psd_db = 10 * np.log10(psd + 1e-20)
```